

DECEMBER 2025

---

# THE INVESTMENT CASE FOR BITTENSOR



**UNSUPERVISED  
CAPITAL**

**SAMI KASSAB  
MANAGING PARTNER**

**SETH BLOOMBERG  
PARTNER**

**RJ GRANT  
PARTNER**

Contact us at [contact@unsupervised.capital](mailto:contact@unsupervised.capital)

WITH CONTRIBUTION FROM



**CRUCIBLE LABS**

# TABLE OF CONTENTS

Introduction	04
<b>Chapter I: Bittensor's Story</b>	<b>06</b>
The Making of Bittensor	07
The Y Combinator Anyone Can Join	09
Our Bittensor Bet	11
<b>Chapter II: Deep Diving Into Subnets</b>	<b>14</b>
The Mechanics of a Subnet	15
Open Source & Building Frontier Tech	18
Types of Subnets	19
Exploring the Ecosystem	22
<b>Chapter III: Valuing TAO</b>	<b>31</b>
How Subnets Capture Value	32
A Valuation Framework for TAO	37
Outlook	43
<b>Disclaimers</b>	<b>44</b>





CRUCIBLE LABS

# All your TAO in one place.

Stake. Allocate. Track PnL. Transfer.  
Everything you need in one Bittensor-native wallet.

[CRUCIBLELABS.COM](https://cruciblelabs.com)

The image shows a smartphone displaying the Crucible Labs mobile application. The top bar indicates the wallet address as "clWallet 5J3mB2...dwal". The main screen features a large white balance display "\$12,450.00" with a green upward arrow and the text "▲ \$2,847.50 (18.42%) 1D". Below the balance, there's a table comparing "Staked" and "Free" amounts: \$9,500.00 and \$2,950.00 respectively. A prominent purple button labeled "Smart Allocate TAO" is centered below the table. At the bottom, a navigation bar includes tabs for "Alpha" and "Activity", along with a menu icon.

Staked	Free
\$9,500.00	\$2,950.00

Smart Allocate TAO

Alpha Activity

Ridges (62)	\$6,010.00
425.00	+\$1580.00 (35.66%)

# Introduction

Generational opportunities often emerge from fundamental market blind spots. Today it's Bittensor.

Over the last two years, Bittensor has quietly built the largest collection of crypto-AI startups, with over a hundred distinct entities operating across agent development, model training, inference provision, and frontier research.

These startups, known as **Subnets**, allow anyone, anywhere, to contribute to a shared vision, all coordinated through an incentive mechanism. This creates a new startup model built for frontier work, giving teams the scale and speed needed to compete with the largest AI labs.

Certain Subnets are already delivering state-of-the-art results, generating millions in revenue, and attracting talent from places like OpenAI and DeepMind to contribute to their missions. They're proving that Subnets can stand alongside traditional AI companies and, in many cases, move faster.

However, the broader market has not yet processed the significance of what's unfolding within Bittensor. This has left TAO, the asset that captures value created across the entire ecosystem, as one of the most undervalued tokens in crypto. **Our base case projects TAO at \$4,800 by December, 2027—a 16x from today's prices. In our bull case, we believe TAO could climb to \$10,800 in that same timeframe.**



Most of the crypto-AI world is still stuck philosophizing about the intersection of the two technologies. The same moral arguments keep resurfacing: *AI centralizes power, crypto democratizes it. AI erodes privacy, crypto protects it. AI captures value, crypto redistributes it.*

These arguments, while well-intentioned, are irrelevant. Virtue and idealism doesn't drive adoption. Superior tech and economics do.

**Crypto-AI wins because it makes AI innovation faster and more efficient—and Bittensor is built on that understanding.**

What crypto does best is combine ownership, economic incentives, and open participation to coordinate human effort at an unprecedented scale. When applied to AI, this coordination unlocks faster iteration, broader contribution, and better resource allocation.

Today, TAO trades at a \$3.0B market cap, while OpenAI raised thirteen times that amount (\$40B) in a single round and is now targeting a one trillion dollar IPO. When you look at what Bittensor has already accomplished, the asymmetric opportunity becomes hard to ignore.

**Just as Bitcoin and Ethereum became the defining networks of their categories, at Unsupervised Capital, we think Bittensor will follow the same trajectory. And we've gone all in on it.**





## The Engine Behind Bitcoin, Applied to AI

Bitcoin started out simply rewarding miners for performing a single task: computing SHA-256 hashes. Simple by design, yet beneath that simplicity, the network's incentives forced a wave of complex innovation.

Miners began competing to improve their efficiency. They designed specialized chips optimized for hashing, engineered advanced data center cooling systems, built custom firmware, and relocated their operations to regions with stranded or excess energy to get cheaper electricity.

Each move was an optimization driven purely by economic incentives to lower costs and increase rewards. Around the clock and across the world, miners competed relentlessly to out-innovate one another.

Because of all their effort, Bitcoin is the largest specialized compute network in the world today, with BTC ranking amongst the ten most valuable assets on the planet.

### **Bitcoin revealed two universal principles:**

- 1. Any process that can be optimized will be, once incentives are attached.**
- 2. A cryptoeconomic protocol can coordinate global resources toward a single objective better than any centralized system could.**

So why shouldn't those same principles be applied to AI? If a cryptoeconomic protocol can drive the optimization of SHA-256 compute, it can just as effectively drive optimization in model training, inference, data collection, or any measurable step in building intelligence.

The mechanism is the same. Incentives coordinate global contributors around an AI objective and give them ownership in the outcome, unlocking a scale and quality of talent that centralized orgs simply can't access. Founders are taking note because in AI, everything comes down to rate of improvement.

And that rate exposes the limits of the traditional startup model. Progress in AI depends on talent density, rapid iteration, and constant experimentation, but startups are constrained by who they can recruit and afford to hire. The closed startup model—where only the people on payroll can contribute—is no match for labs operating at nation-state scale budgets with the ability to hire anyone they want.

Bittensor is designed around this reality. It replaces the closed startup model with one where anyone can contribute and be rewarded for improving the system. Instead of raising enough capital to compete with national-scale labs, founders design incentives that attract global resources, talent, and experimentation. The systems that win will be the ones that learn and improve the fastest. Crypto gives AI startups the coordination engine needed to reach that pace.

## CHAPTER I

# Bittensor's Story

- The Making of Bittensor
- The Y Combinator Anyone Can Join
- Our Bittensor Bet



## The Making of Bittensor

### Bittensor Timeline



Bittensor's story begins in January 2021, when it launched as a simple network inspired by Bitcoin. It shared many of the same design principles: a fixed supply of 21 million tokens, a fair launch, and a programmatic emission schedule that distributed new TAO every day for computational work.

In its early years, Bittensor focused entirely on the single task of serving inference. Miners processed text prompts through language models and competed to return the fastest, most accurate, and lowest-cost responses. The system worked—and it worked well.

But as the network proved itself, more questions emerged. If the network could incentivize one AI task, why stop there? Why shouldn't Bittensor expand to many tasks? And what if anyone could build one of these task-focused networks within Bittensor?

Those ideas came to life in October 2023 with the Revolution upgrade. It introduced Subnets, letting independent teams launch their own specialized networks within Bittensor, each with its own goal, contributors, and community. TAO was distributed to them based on the value they delivered back to the network.

But as Subnets scaled from a handful to dozens, distributing TAO to each Subnet became a bottleneck. The initial approach involved large stakeholders manually determining TAO distributions after subjectively evaluating each Subnet. That worked to bootstrap the system, but it prevented the network from scaling and gave a small group discretionary control over billions in capital allocation.

In February 2025, the Dynamic TAO upgrade introduced a fix to those problems using a novel market-based emission allocation system. Dynamic TAO gave every Subnet its own token and allowed anyone to direct TAO emissions by staking their TAO into the Subnets they believe in. The more stake a Subnet attracts, the larger its share of daily TAO emissions. This turned Subnet tokens into live market signals for where TAO should flow, replacing the discretion of a small group. It was the most important upgrade in Bittensor's history.

Bittensor has evolved a lot since inception, but the mission has been consistent: build the most powerful coordination and capital allocation system for AI development. Four years later, that mechanism now directs billions in resources across 100+ specialized networks.



## The Tensorpunks

Bittensor's success today can be traced back to its self-organizing community. Most projects never manage to build one since they can't be manufactured with a large treasury, VC connections, or grants to mercenary devs. A real community emerges from people who choose to show up, contribute, and stay because they believe in the mission.

**Bittensor's community is global, grassroots, and relentless. Two deliberate choices made this possible:**

- 1. A Bitcoin-like fair launch, which made it attractive for early contributors to invest and participate; and**
- 2. The decentralized AI values held by Bittensor's founders, Jacob Steeves and Ala Shaabana.**

The first point is underappreciated. Lasting networks need a community that actually cares and is committed for the long haul. The only reliable way to create that is to make your early believers win (rich). The most direct way to do that is through early participation, where normal people, not just VCs and institutional investors, get the chance to capture substantial upside. This is best executed via a fair launch mechanism.

When you do that, you attract owners and create zealots who become bonded to the network both economically and emotionally. They're incentivized to see it keep winning, and with the capital they earn, they reinvest back into the ecosystem.

This is why Bittensor has one of the most impressive grassroots ecosystems in crypto. It has the same early-days energy seen in Bitcoin and Ethereum: a bottom-up movement aligned around a clear mission —building decentralized AI—with the community willing to stick with it for the long haul.

What makes this more notable is how little top-down support exists. The OpenTensor Foundation doesn't have a business development team, a marketing lead, or a war chest from a pre-mine to fund operations. It is intentionally lean and shrinking as the network decentralizes, with the goal of eventual dissolution.

Almost everything you see around Bittensor today was built by the community. The first exchange listings for TAO were coordinated by people in a Telegram chat with no formal connection to the foundation. The community members made wealthy by the chance to invest early were the ones who funded the initial Subnet teams, built early infrastructure, paid for marketing, and created the first interfaces to the network.

Around Bittensor, a loose network of contributors has grown into independent professional orgs handling everything from public infrastructure to incubating to protocol development. Crucible Labs, Yuma, Taostats, Latent Holdings, Manifold Labs, Macrocosmos, Tensors, Covenant, GTV, Bitstarter, Stillcore and yours truly are just a few of the teams now driving the network forward alongside the foundation.



## The Y Combinator Anyone Can Join

Bittensor operates like an incubator on autopilot. It's distributing over a billion dollars worth of TAO emissions each year (as of Nov. 2025) to the best-performing teams building within it. And as TAO's price grows, so does the scale of that funding.

To understand what Bittensor enables, consider one of the most successful startup accelerators, Y Combinator (YC). Since 2005, it's been Silicon Valley's most prestigious startup accelerator—funding the likes of Airbnb, Stripe, Coinbase, and Reddit. YC's model is to invest \$500K in promising teams for 7% equity, provide mentorship for 3 months, then showcase them to investors at Demo Day. Bittensor took this model, removed every gatekeeper, and improved it.

	<b>Y Combinator</b>	<b>Bittensor</b>
<b>Application</b>	Write essays, get references, hope you know someone	Pay the network registration cost and start competing immediately
<b>Selection</b>	Partners in a closed room debate your potential	The open market determines your potential
<b>Funding</b>	\$500K once, find your own investors after	Daily TAO rewards determined by the market
<b>Capital Asset</b>	Equity: Aligns founders and venture investors, dilutes every new round	Tokens: Aligns founders, investors, and all other contributors
<b>Timeline</b>	3 months then Demo Day	Infinite—compete as long as you can
<b>Competition</b>	Against 200-260 other startups in your batch	Against everyone, everywhere, always
<b>Mentorship</b>	Weekly dinners with Paul Graham	From thousands within the ecosystem
<b>Liquidity</b>	IPO or acquisition (maybe, in 5-10 years)	Public market for the token on day-one



## TAO as Exposure to the Entire Portfolio

Even the best accelerators have trouble picking winners. Humans are terrible at predicting which weird ideas will become billion-dollar companies. It's one of the core reasons why Bittensor shifted from its original system (a small group of stakeholders that decided on emission allocation) to a market-based mechanism. Even the best VCs are wrong 80% of the time.

Instead of a handful of VCs making decisions, Bittensor lets the market make the decision on how much funding the project deserves. Teams need to quickly build real utility and tell a compelling story to continue getting funding from the network. Those unable to do so see their funding fade away. This creates pure meritocracy without politics, bias, or human error.

And this creates something traditional accelerators can't match. YC can only review a finite number of applications, run a few batches a year, and support a limited set of founders. Bittensor can run an infinite number of experiments simultaneously.

Viewed this way, Bittensor behaves less like a single project and more like a portfolio of them, all accessible through TAO. It doesn't need every Subnet to work. As with any venture portfolio, a few breakout successes are enough to make the entire system a generational outcome.





## Our Bittensor Bet

As an investment fund, we've identified three pillars that define our thesis:

- 1. Bittensor's architecture is the best environment for crypto-AI builders;**
- 2. The smartest AI developers in the world will keep showing up to mine Subnets; and**
- 3. Subnet successes will compound, creating network effects that pull the best builders into Bittensor**

If these three pillars continue to hold, we anticipate TAO's value to continue climbing.

---

### 1—Bittensor is Superior For AI Builders

Blockchains were never designed with AI in mind. They were built for trustless record-keeping in binary environments. Deterministic computation. Every node must agree on the exact same answer. A transaction is valid or invalid. A smart contract's conditions are either met or not.

But AI does not work in absolutes. Models are probabilistic by nature. Identical prompts put through the same model can return different responses. Traditional consensus mechanisms cannot handle that, because there is no single "correct" answer for every node to agree on. If Ethereum validators were constantly seeing different balances for the same address, the chain of record would break.

On top of that, you could never run real AI workloads onchain. They are far too computationally intensive. So AI needs a consensus mechanism that can work with its probabilistic outputs and its heavy compute requirements. Bittensor solved this with [Yuma Consensus](#).

Yuma Consensus made it possible for a decentralized network to agree on quality, not just binary outcomes. Validators can come to consensus on subjective questions like: How 'good' is this miner's model output? How accurate is this prediction? How fast is this miner's inference? Crucially, all mining and validation activity occurs offchain, with only the validators' final scores of the miners being recorded onchain.

The bottom line is Bittensor is the only Layer-1 blockchain designed specifically to work with AI.



## 2—The Deepest Pool of Mining Experts

Bittensor doesn't exist without the miners. They're the ones showing up every day to do the actual work —training vision models, hosting the latest and greatest open-source models, optimizing image generation pipelines, etc. If the miners stop optimizing for the incentive mechanism, competition stalls. And if competition stalls, Subnets stop improving.

Miners are Bittensor's most powerful network effect, and they're the reason founders want to build here. The network has attracted an unusually specialized mix of talent ranging from PhDs from Meta competing against teenagers who taught themselves machine learning on YouTube. This is part of what makes Bittensor so powerful. It's become an aggregator of specialized intelligence on the supply side.

### **Within the 128 Subnets, you'll find:**

- **PhDs training computer vision models;**
- **Young technical prodigies optimizing inference libraries;**
- **OpenAI and DeepMind engineers moonlighting to build software-engineering agents;**
- **Professional mining firms scanning the network for gaps they can fill;**
- **Operators running large data centers providing GPUs; and**
- **Seasoned professionals participating to earn supplemental income.**

They all meet here, competing and collaborating in the same open economy. And when a new Subnet launches, they're ready to plug in from day one if their skills fit the scope. Launching a new Subnet on Bittensor is like launching a DeFi protocol where liquidity is already at its deepest—you get to launch into an active, liquid market immediately.



## 3—A Positive Sum Ecosystem

Bittensor is not a zero-sum game like most ecosystems, where teams compete for the same users or liquidity. Here, they're part of the same economy. The more Subnets that succeed, the easier it becomes for new ones to bootstrap.

Launching a Subnet is different from launching a standalone network. When one Subnet gains traction, demand for TAO increases because it's needed to acquire each Subnet's token. As TAO appreciates, the entire network benefits:

- **Subnet tokens rise in value** - Since all Subnet tokens are priced in TAO, appreciation lifts the entire ecosystem.
- **Miner quality improves** - Miner reward pools increase in value alongside TAO, attracting more resources and talent.
- **Ambitious projects become viable** - Higher TAO prices make it easier for new teams to launch and fund capital-intensive experiments like trillion-parameter model training.

What that means in practice is that the success of one Subnet compounds across the entire network.

**We're watching two metrics: the numbers of both quality Subnet founders and miners joining. If those keep trending up and to the right, the investment thesis for Bittensor remains strong.**

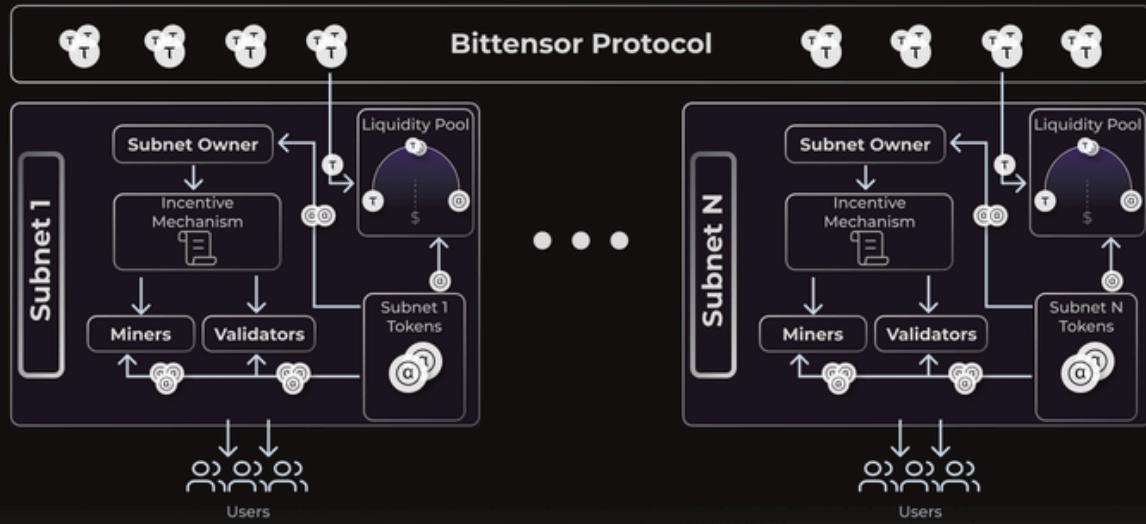
## CHAPTER II

# Deep Diving Into Subnets

- The Mechanics of a Subnet
- Open Source & Building Frontier Tech
- Types of Subnets
- Exploring the Subnet Ecosystem

## What is a Subnet

### Bittensor Subnet Architecture



Subnets are to Bittensor what smart contracts are to Ethereum. They serve as the core building block for creating applications on the network.

At its core, however, a Subnet operates like a decentralized startup, with the Subnet Owner acting as the startup founder. The Subnet Owner defines the main properties of the Subnet, including the vision and the tasks that need to be executed. These properties are represented with an incentive mechanism, which serves as the objective and rules of the network.

Just as every startup has employees, Subnets have miners. The set of miners on a Subnet compete to build the best solution based on the incentive mechanism. Validators evaluate miners using a set of key performance indicators specified by the Subnet Owner. For a Subnet, the miners become a globally distributed and competitive workforce.

For example, consider a Subnet with the objective of developing an inference-as-a-service startup:

- **Subnet Owner:** Designs the grading rubric that defines how miners can earn the highest incentives —such as maintaining high inference accuracy, low latency, and consistent uptimes.
- **Miners:** Host models and optimize their inference systems to be the best at processing user prompts. Better performance leads to higher payouts for a miner.
- **Validators:** Evaluate each miner according to the rubric and assign scores that determine miner rewards.



While the vast majority of Subnets today are AI-focused, it's not a requirement. Any measurable service or task that can be rewarded based on performance can be codified as a Subnet.

The beauty of the Subnet design is how easy it is to launch one. All that's required is an incentive mechanism. Bittensor handles the rest:

- Tokenomics are pre-defined,
- Decentralized exchanges are integrated into each Subnet with liquidity provided by the chain; and
- No blockchain-specific infrastructure knowledge is required, as each Subnet inherits a fixed ruleset that builders don't need to recreate.

By removing crypto's steep learning curve and making the launch process simple, Bittensor pulls in top-tier talent from traditional AI and other legacy fields. It has turned the network into a launchpad where anyone with an idea and an incentive mechanism can easily spin up a new crypto-powered startup.

While launching a Subnet is quite easy, surviving is tough. Bittensor is currently limited to 128 Subnet slots, and every one of them is filled. This enforced scarcity drives up the quality bar for participants. Over time, this cap is expected to increase gradually, with the long-term goal of supporting an effectively unbounded number of Subnets. When a new team wants to launch, the lowest-performing Subnet is deregistered. This creates constant competition not only to avoid the bottom, but to climb the top of the leaderboard, since higher-ranked Subnets earn a larger share of TAO emissions.

Each Subnet contains its own decentralized exchange (DEX) implemented as an automated market maker (AMM) liquidity pool. When Bittensor distributes emissions, it does so by injecting TAO into one side of the pool and an equalizing amount of Subnet tokens into the other side, maintaining pool balance. These pools serve as the primary venue for contributors to convert their Subnet tokens back into TAO, and for investors to gain exposure to Subnet tokens.

The protocol then observes net trading volume into each Subnet's AMM pools and uses those signals to determine how TAO emissions should be allocated across Subnets. **Ultimately, Subnets with stronger market interest receive more TAO emissions, while the Subnet with the lowest market interest is automatically removed when a new Subnet registers.**



Subnet	Emission	Market Cap
1. Chutes SN 64	6.7%	\$82M
2. Lium SN 51	4.1%	\$50M
3. Targon SN 4	4.4%	\$42M
4. Ridges SN 62	1.5%	\$42M
5. Vanta SN 8	3.6%	\$35M
6. Affine SN 120	6.7%	\$30M
7. Templar SN 3	2.5%	\$24M
8. Gradients SN 56	1.6%	\$24M
9. IOTA SN 9	2.2%	\$19M
10. Hone SN 5	1.0%	\$17M

Data as of 12/05/25

# Market Map



## Compute & Infrastructure

	SN 64	Chutes Serverless AI Compute
	SN 4	Targon Confidential Cloud Compute
	SN 51	Lium GPU Rental Marketplace
	SN 39	Basilica Trustless GPU Marketplace
	SN 12	Compute Horde GPU Rental Marketplace

	SN 95	Actual Computer Consumer Hardware Inference Network
	SN 27	NodeXO GPU Compute Marketplace
	SN 2	Subnet 2 ZK Proofs
	SN 11	Dippy Media Generation Platform
	SN 48	Quantum Compute Advancing Quantum Compute Research

	SN 63	Quantum Innovate Advancing Quantum Compute Research
	SN 75	Hippius Decentralized Storage
	SN 85	Vidao Video Processing
	SN 14	TAO Hash Bitcoin Mining and Hashrate Marketplace
	SN 89	InfiniteHash Bitcoin Mining

## Robotics

	SN 26	Kinitro Incentivized Embodied Intelligence
	SN 47	Reboot Environment for Robotics Development
	SN 124	Swarm Global Arena for Autonomous Drone Training

## Data Pipeline

	SN 13	Data Universe Data Scraping and Storage
	SN 33	ReadyAI Data Annotation and Structuring
	SN 54	Yanez Repository and Compliance Database Generation
	SN 37	Aurelius Alignment Dataset Generation
	SN 71	Leadpoet Marketing and Sales Dataset Generation
	SN 52	Dojo Human Generated Data and Labeling
	SN 22	Desearch Search and Data Aggregation
	SN 96	Flock Off Fine-tuning Dataset Generation
	SN 42	Gopher Data Scraping
	SN 46	Resi Real Estate Data Aggregation

## DeFi

	SN 35	Cartha Perpetuals DEX
	SN 106	VoidAI Liquidity Provisioning and Bridging
	SN 10	Swap Core DeFi Primitives
	SN 15	BitQuant On-chain Quantitative Analysis
	SN 67	Tenex Spot Margin Protocol

## DeSci

	SN 68	Nova Drug Screening and Discovery
	SN 18	Zeus Environmental Forecasting
	SN 57	Gaia Geospatial Machine Learning

## Misc

	SN 93	Bitcast Creator Marketing
--	-------	------------------------------

## Decentralized Training

	SN 81	Grail Verifiable Inference and Post-Training
	SN 9	IOTA Model Parallel Pretraining
	SN 3	Templar Large Scale Pretraining

## Predictive Systems

	SN 41	Sportensor Sports Prediction
	SN 30	Bettensor Sports Prediction
	SN 123	Mantis Bitcoin Price Prediction
	SN 31	Candles Crypto Asset Price Prediction
	SN 53	Efficient Frontier Institutional-Grade Trading Strategies
	SN 55	Precog Bitcoin Trading Strategies
	SN 6	Numinous Event Prediction
	SN 88	Investing Optimized Staking Strategies
	SN 8	Vanta Long/Short Trading Signals
	SN 50	Synth Synthetic Price Path Data
	SN 79	Taos Simulation of Automated Trading

## Model & Agent Development

	SN 62	Ridges Software Engineering Agents
	SN 60	Bitsec Software Auditing Agents
	SN 61	Red Team Cybersecurity
	SN 34	BitMind Deepfake Image and Video Detection
	SN 32	ItsAI AI Text Detection
	SN 17	404-GEN 3D Generative Model Development
	SN 36	Autoppia Computer-Using Agents
	SN 44	Score Visual Intelligence and Computer Vision
	SN 72	StreetVision Object and Image Classification
	SN 43	Graphite General Algorithmic Optimization
	SN 1	Apex Algorithmic Optimization

Subnets as of December 2025

Not all subnets included



## Open-Source as The Best Way to Build Frontier Tech

Open-source is the quiet engine that built the modern world. The internet runs on open protocols like TCP/IP and DNS. Linux is the backbone of modern computing. Android defined the modern smartphone era. PyTorch and TensorFlow launched the AI boom.

These technologies succeeded because of the “open collaboration” element of open-source. Closed systems in the long-run, no matter how well-resourced, can’t coordinate enough expertise at the speed and scale required at the frontier. Open-source lets anyone, anywhere, contribute, creating a kind of collective intelligence. A system where the world’s creativity, expertise, and obscure edge-cases compound in one place.

From a first-principles view, open-source should be the best way to build startups working on frontier technologies. It gives you speed, reach, and the sharpest minds in the industry. But it hasn’t worked well as a startup model because: (1) startups need to capture value to survive and (2) contributors need a path to upside, which open-source has previously never offered.

**Subnets change this. They’re a new kind of startup built on open-source participation, with a native mechanism for contributors to raise funding and capture value.**

## Subnets = Open-Source Start-Ups

Every Subnet has a native token that represents ownership, coordinates contribution, and funds the system. And since the token trades immediately, investors can fund the network while contributors receive liquid upside for their work on day one.

Linux created billions in value, but the developers who built the kernel captured almost none of it. IBM paid \$34 billion for Red Hat, not to the early contributors who made Linux possible. Subnets fix this by introducing an ownership layer natively, letting builders collaborate openly while sharing in the value they create.

## Types of Subnets

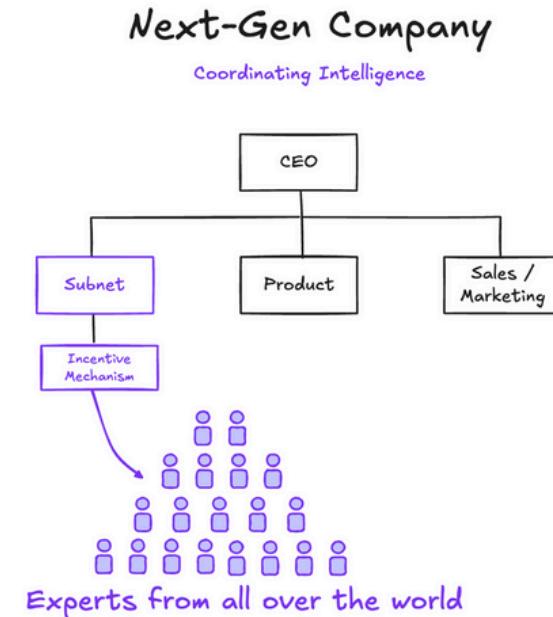
There are two main types of Subnets, those that:

- 1. Coordinate a distributed workforce—researchers, engineers, subject matter experts; and**
- 2. Coordinate digital resources like compute.**

For both models, open collaboration is their strategic advantage.



# Subnets That Coordinate Intelligence



Take a Subnet focused on developing software-engineering agents. The founder designs one thing: an incentive mechanism that rewards whoever produces the top-performing agent. That's it. Once that objective function is live, anyone—an engineer at a top lab or a self taught developer—can show up, compete, and earn purely on merit. The efficiency of the system comes from encoding performance standards directly into software. Designing these mechanisms is hard. Subnet Owners iterate constantly on them. But once deployed, the efficiency they unlock is unmatched.

The founder's job shifts from “building the agent” to turning miner outputs into a product. They take the agents the network produces and handle productization—building apps, user interfaces, establishing distribution channels, and managing partnerships.

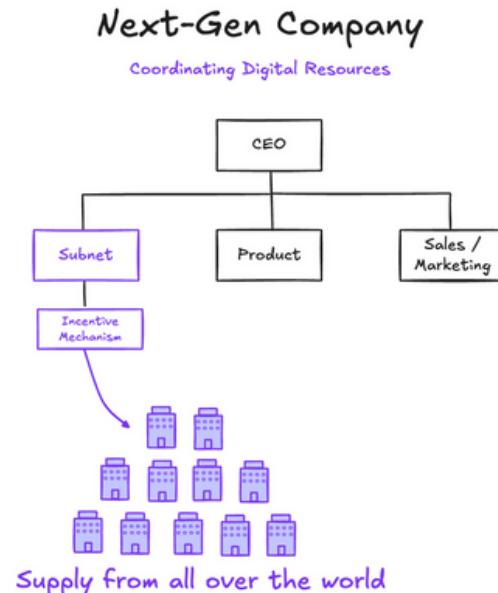
This flips the traditional startup model. Instead of building a headcount-heavy organization, Subnet founders design incentive systems that attract a global contributor base that become the Subnet's distributed R&D team. Participation is open and programmatically verified, so trust isn't required. That becomes the Subnet's competitive edge. Domain experts can show up, plug in, and get paid. No interviews, no coding tests, no performance reviews.

The organizational overhead shrinks dramatically. The incentive mechanism handles talent selection and performance management automatically, allowing the system to scale. High-performers rise. Low-performers get pushed out.

This Subnet model unlocks an entirely new way to build companies in AI. In our view, it may be the only viable path for startups to compete with the incumbents since AI's feedback loops keep reinforcing their advantage. The only realistic counterweight is coordination at scale and Subnets make that possible. **When rate of innovation is the only moat that matters, open participation becomes the ultimate advantage.**



# Subnets That Coordinate Digital Resources



Now, take a Subnet building a neocloud. The founder designs an incentive mechanism that rewards compute providers based on how much compute they offer, the type of compute, and how reliably they stay online. Any compute provider can participate—a small-to-medium-sized data center, a university lab, or an individual GPU owner.

The incentive mechanism handles the trust problem automatically. It verifies that providers have the hardware they claim, confirms workloads are being executed honestly, and filters out unreliable or low-quality providers. The founder doesn't have to deal with any negotiating, contracts, or manual onboarding. Supply becomes permissionless, allowing for the fair market value of compute to surface.

This frees the founder to focus entirely on the product layer—the API, developer experience, distribution, enterprise integrations, and go-to-market—instead of spending their time sourcing, managing, and policing compute providers.

A Subnet coordinating digital resources naturally becomes a two-sided market. And these markets can get extremely specialized, because the incentive mechanism can be tuned to reward one very specific form of supply. Specialization becomes the advantage, making the market more efficient, cheaper to operate, and better aligned with what customers actually need.

**Bittensor enables liquid compute markets.** Providers can switch between Subnets with near-zero friction. No contracts or negotiations, just redirect compute and start earning. When a training Subnet shows strong demand, it attracts new supply. When an inference Subnet's performance slips, providers reallocate. Compute flows dynamically to its highest-value use, directed by the market. As more compute owners look to monetize idle capacity the way energy producers monetize excess power, the need for a liquid, on-demand marketplace becomes obvious. And Bittensor can support it.



## Subnet Native Funding

Subnets start out like any other startup, they need funding before they can generate revenue. In the traditional world, founders raise venture capital through equity to bootstrap operations until the product works and paying customers show up. In Bittensor, funding is achieved via token emissions.

Every Subnet launches with its own token and a fixed emission schedule designed to fund the system from day one. Those emissions are split across contributors: 18% to the Subnet Owner, 41% to miners, and 41% to validators. This becomes the Subnet's built-in financing mechanism.

Subnets receive TAO emissions from the network based on the value they create relative to others. That TAO goes into the native liquidity pool, paired with the Subnet token. This mechanism effectively backs each Subnet token with TAO liquidity. The pool is how contributors realize their rewards. It's also how investors and speculators gain exposure.

In effect, Subnet token holders and TAO holders become the early-stage investors, providing funding for the Subnet during its bootstrapping phase. Their capital supports miners, funds R&D, and gives the Subnet time to reach product-market fit.

The goal for every Subnet is to eventually become sustainable, meaning to generate enough value and attract enough demand that emissions are no longer the main source of funding. We explore that transition in Chapter 3.



## Exploring the Subnet Ecosystem

With 128 Subnets live—and new ones joining every week—the ecosystem is evolving faster than anyone expected. Due to the Darwinian nature of the ecosystem, teams are pushing the boundaries of what this organizational model can do, trying ideas that would never be feasible inside a traditional company or incubator.

Amid all that experimentation, a handful of Subnets have already shown what this architecture is capable of and the direction the ecosystem is likely to continue expanding in.

- **Ridges** shows how a Subnet can bootstrap world-class AI systems and mine human intelligence at scale.
- **Chutes** proves a Subnet can power a real product at massive scale and go head-to-head with billion-dollar incumbents.
- **404-GEN** demonstrates how global talent can be coordinated to accelerate R&D far beyond what any in-house team could achieve.
- **Targon** shows how a Subnet can serve enterprise customers at scale.
- **IOTA** shows that Subnets can take on deep, technical frontier problems that were previously reserved for well-funded labs.
- **BitMind** shows how a Subnet can produce state-of-the-art, domain-specific models.
- **Synth** proves that Subnets can thrive outside of core AI markets entirely.



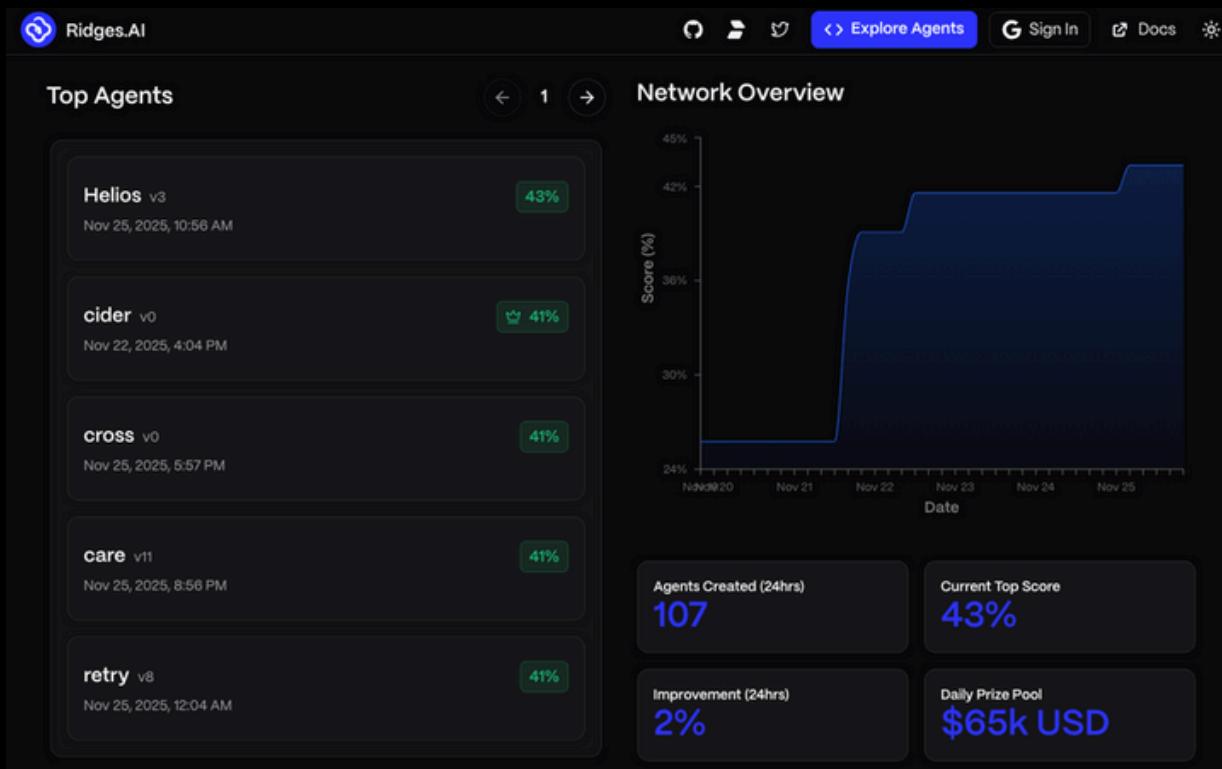
## Ridges

Ridges turned software-agent development into an open tournament, rewarding contributors for building the strongest coding agents. After just four months, Ridges produced the top-ranked, fully open-source software engineering agent on the SWE-Bench leaderboard. It even outperformed Anthropic's Claude 4 with tool support.

And it did so with minimal internal resources: a lean team of Shakeel (the Subnet Owner) + three Waterloo undergrad interns, no venture funding, and not a single line of agent code written by the core team. Instead of building in-house, Ridges leaned on the prize pool incentives—roughly \$35K per day at recent token prices—to attract contributors from OpenAI, DeepMind, and hundreds of independent developers, all competing to build the best coding agent.

On the Subnet, miners submit agents as single Python files that leverage open-weight AI models. Every agent submission by miners must be fully open-sourced, which is one of the main reasons the Subnet has moved so fast. Open-sourcing agents speeds up iteration, eliminates duplicate work, and compounds progress.

After proving that the system could produce state-of-the-art results, Ridges is now moving beyond research and benchmarks. The team has developed an agent platform where engineers or even non-technical users can create and assign tasks for agents to complete.

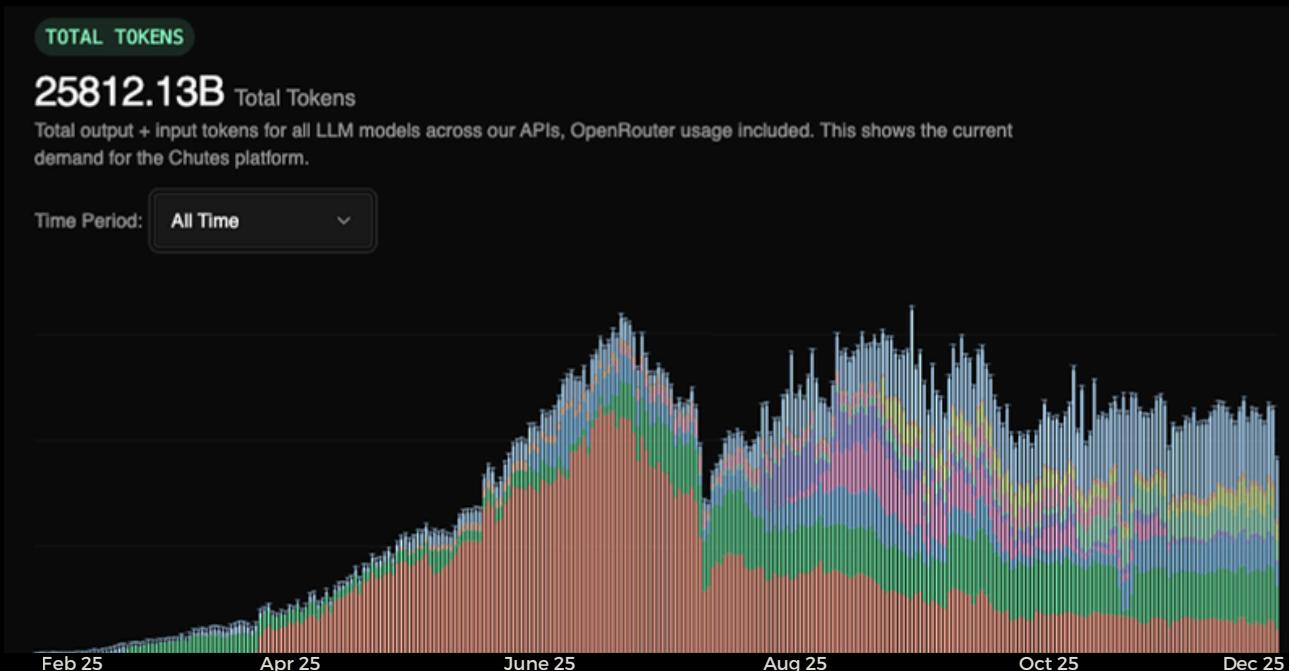
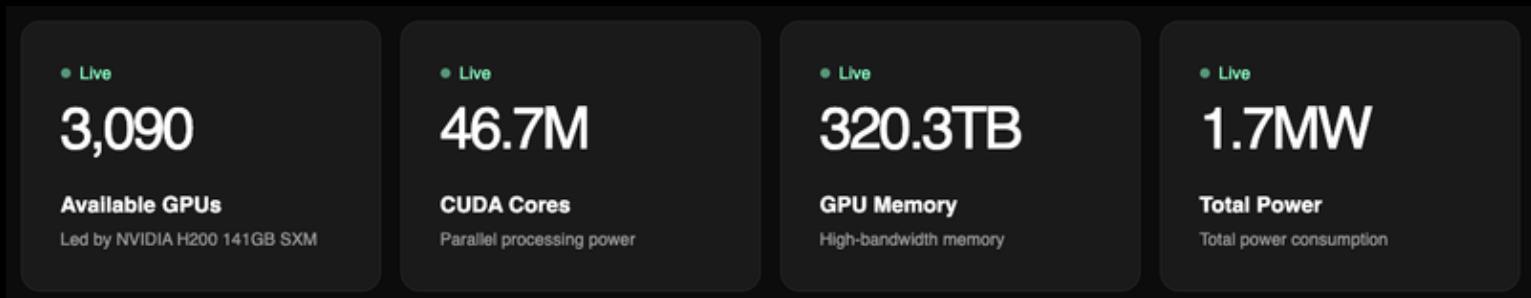




## Chutes

**Chutes** is serving more than **110 billion** tokens per day and ranks among the leading providers of open-source models on OpenRouter, with roughly a quarter of its daily volume coming from that platform alone.

It delivers a production-grade interface where users can purchase competitively priced subscriptions for serverless inference, all without ever knowing Bittensor is under the hood. Behind the scenes, a distributed network of miners—mostly small and mid-sized data centers—supplies the GPUs that power its models. Many of their miners run data centers with idle GPUs and are willing to contribute them to the Subnet at prices far below market rates. Chutes built its platform so providers can reclaim their GPUs at any time without interrupting end-user workloads. This cost advantage lets Chutes offer cheaper inference to companies with large AI bills, winning customers it otherwise could not reach.



Data as of November 26, 2025

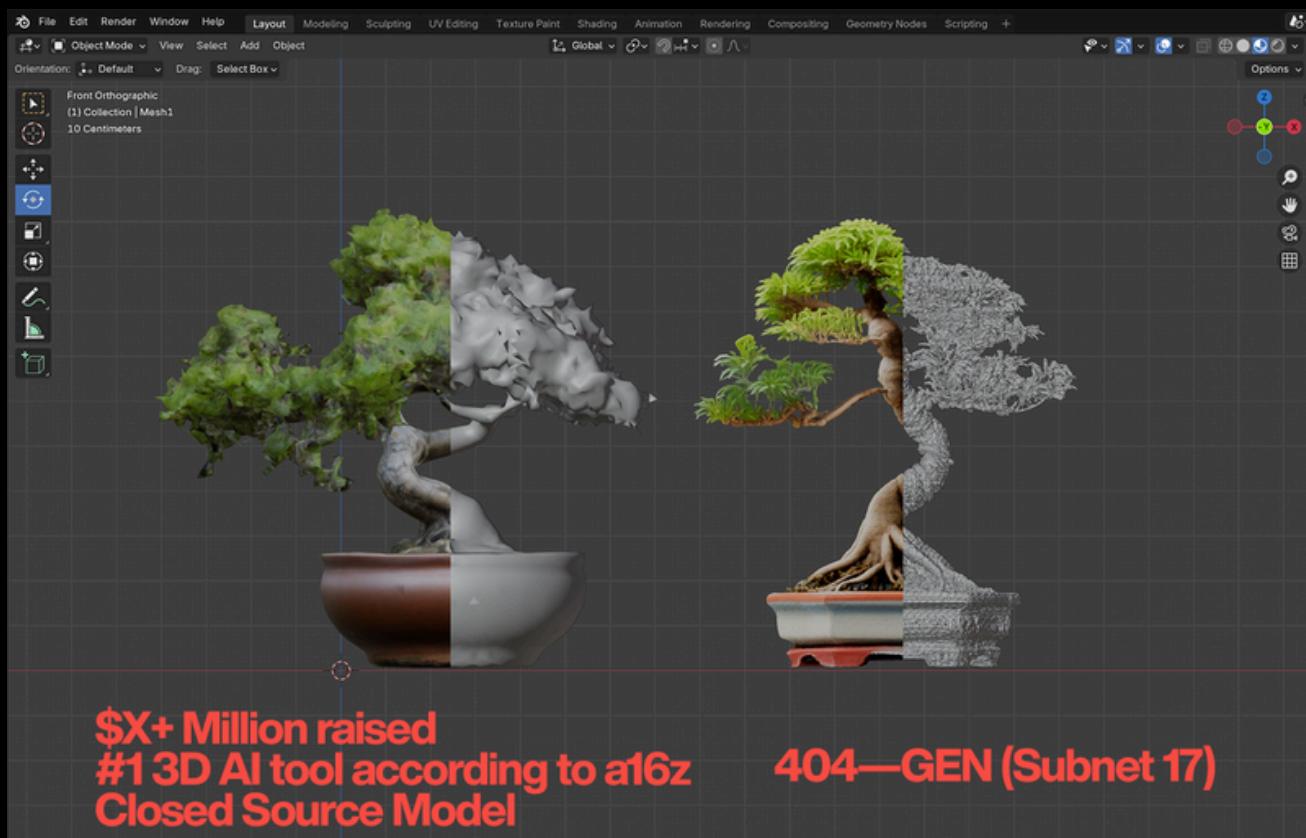


## 404-GEN

Before discovering Bittensor, Ben James, founder of **404-GEN**, was already building Atlas, an end-to-end design platform for AAA game studios. A core piece of Atlas requires best-in-class 3D asset generation. Instead of opting for hiring an expensive R&D team to build and maintain generative models internally, Ben launched the 404-GEN Subnet to let the world compete as his R&D team.

Researchers compete in an open tournament to build the best 3D generation models, submitting their work for testing against industry benchmarks and head-to-head duels with other miners.

And the system has been working well. 404-GEN's models rank highest in aesthetic quality compared to Meshy, CSM, Hunyuan, and Trellis, making it one of the most advanced 3D model generators in existence.





## Targon

Targon is a cloud provider focused on end-to-end verifiability and privacy through hardware-enforced security. Acting as the buyer of last resort for data centers globally, many of whom are sitting on massive amounts of idle capacity, Targon has aggregated this supply via a traditional order book mechanism, securing over 1,600 H200 GPUs. It is a win-win: datacenters operate at 100% utilization and Targon customers enjoy the benefits of a structural cost advantage.

Targon's secret sauce is their Targon Virtual Machine (TVM). It uses hardware-backed confidential computing across NVIDIA, Intel, and AMD chips to guarantee security cryptographically. This minimizes trust requirements, enabling enterprises to move sensitive workloads to Targon with the certainty that their data remains fully encrypted and verifiable.

The team has built [Targon.com](#), centered around confidential cloud compute, and is working on new apps on top of the Subnet including a search app and chat app. As more apps launch, demand for Targon compute grows, which attracts more supply, lowering costs and accelerating the flywheel.

**TARGON**

**Inventory**

[Create Rental](#) [Add SSH Key](#) [Settings](#)

<b>1x NVIDIA H200 Small</b> from <b>\$2.00/hr</b> Run fully secure NVIDIA H200s on TVM. <a href="#">Learn more</a> <a href="#">Deploy</a>	<b>2x NVIDIA H200 Medium</b> from <b>\$4.00/hr</b> Run fully secure NVIDIA H200s on TVM. <a href="#">Learn more</a> <a href="#">Deploy</a>	<b>4x NVIDIA H200 Large</b> from <b>\$8.00/hr</b> Run fully secure NVIDIA H200s on TVM. <a href="#">Learn more</a> <a href="#">Deploy</a>
--	---	--

**ALL AVAILABLE GPUS**

<b>NVIDIA H200 - Small</b> Run fully secure NVIDIA H200s on TVM.	Price <b>\$2.00/hr</b>	175 GB RAM	1x NVIDIA-H200	<a href="#">Deploy</a>
<b>NVIDIA H200 - Medium</b> Run fully secure NVIDIA H200s on TVM.	Price <b>\$4.00/hr</b>	350 GB RAM	2x NVIDIA-H200	<a href="#">Deploy</a>



## IOTA

**IOTA** is tackling one of the hardest technical problems in AI: decentralized model training using heterogeneous hardware scattered around the world.

Although still early, IOTA has demonstrated that their pipeline parallelism training approach works and is now focused on making the system more efficient and scalable. Their latest upgrade allows anyone with a MacBook to contribute to training, essentially the next evolution of SETI@home or Folding@home, except now it's *training@home*.

Decentralized training working at scale would decouple model training from governments and the handful of companies that control most of the world's compute, and it would broaden who can participate in building frontier AI. Asymmetric opportunities like that don't come around often.





## BitMind

**BitMind's** mission is to build state-of-the-art consumer and enterprise products for AI content detection. Their Subnet incentivizes miners to train classifier models that accurately determine whether an image or video is real or AI-generated.

The Subnet splits its rewards between the AI models that successfully evade detection and the models that successfully determine real vs. fake content. This adversarial architecture pushes both the quality of fakes and the quality of detection models higher, making sure Bitmind is always at the bleeding edge. The best-performing miner model is then used as the foundation for BitMind's products.

This architecture has incentivized models that have outperformed competitors across industry-standard benchmarks. BitMind's consumer apps leveraging those models have reached over 80,000 monthly active users, and their enterprise APIs handle more than 2.6 million requests per week.

BitMind API

```
curl -X POST "https://api.bitmind.ai/v1/34/detect-image"
-H "Authorization: Bearer {token}"
-H "Content-Type: application/json"
-d '{
  "image": "https://picsum.photos/"
}'
```

AI Analysis Result

AI GENERATED  
99.6% Confidence

Confidence Level  
Very High Confidence

BitMind

3:56

Plus 8:40

Detect AI generation

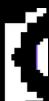
Create with AI

How to import from social media?

More

Keep your skills sharp with our game  
Test your AI detection skills and earn points to compete on the Leaderboard.

Home History Score



## Synth

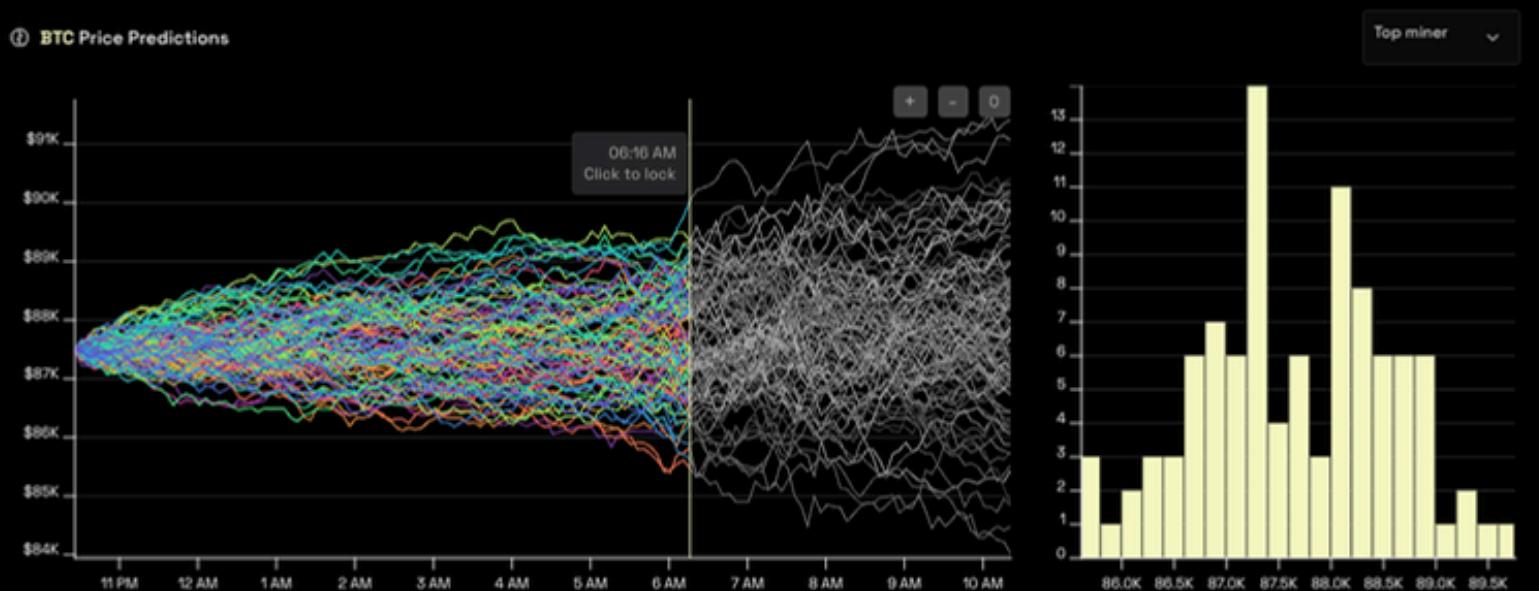
Not all Subnets are building AI products, even if the influence of AI is felt throughout the stack.

**Synth** is building a high-frequency trading firm and market-making services on top of the Subnet. Quant researchers compete to submit the best price forecasting models to the Subnet, providing Synth with over 12B data points daily that feed their metamodel. In the first two months of live trading, Synth is up ~2100% on the Up/Down crypto markets on Polymarket.

Having demonstrated success in its proof-of-concept, Synth is scaling up its trading activities, market making, and asset listings, with each revenue stream funding more token buybacks. Synth is also expanding into new prediction markets and platforms, adding new assets for miners to forecast, and growing volume traded by entering crypto & equity options markets.

The Subnet token sits at the center of Synth's competitive edge. As the token value grows, Synth asks more of its miners, increasing the speed and frequency of their price forecasts, as well as the number of assets included. Synth targets a consistent stream of buybacks given the revenue is correlated with volatility, not market prices or crypto activity.

## MINER PREDICTIONS





## A Decentralized AI Lab

From the beginning, Bittensor's focus has been on figuring out how to fund, develop, and deploy intelligence outside the structure of centralized AI labs. That vision is still alive today, and it's becoming more real every month. The current ecosystem looks like a decentralized AI lab assembling the full stack needed to build intelligence—from data collection (Dataverse), to storage (Hippius), to compute (Targon, Lium), to pre-training (IOTA, Templar), to reinforcement learning (Affine), to inference (Chutes).

Within that stack, some Subnets operate like research groups pushing the frontier—decentralized pre-training, large-scale RL, new training paradigms—while others function like service teams, providing value today and generating revenue. The ecosystem needs both for the model to sustain itself.

This balance still matters as Bittensor expands beyond AI. Whether a Subnet trains models or builds creator-marketing tools, the network depends on having both revenue-generating teams and high-risk research efforts.

Google's 70/20/10 rule offers a useful analogy. Most of its resources went toward advancing the core business, some toward adjacent opportunities, and a smaller share toward high-risk innovation. We think Bittensor's ecosystem should follow a similar pattern. Roughly 70 percent of Subnets should focus on becoming cash-flow positive—earning more revenue than they spend on miner emissions—while the remaining 30 percent can pursue high-risk, high-reward research. The revenue-generating Subnets effectively subsidize the research Subnets, creating a balanced, self-funding ecosystem.

## CHAPTER III

# Valuing TAO & Its Ecosystem of Subnets

- How Subnets Capture Value
- Valuing Specific Subnets
- A Valuation Model for TAO



## How Subnets Capture Value

Valuing a Subnet is the natural starting point for valuing TAO itself.

Subnets function as two-sided markets. Miners compete to produce the best outputs, and the Subnet Owner builds a product and go-to-market around those outputs. Neither can succeed without the other, and both are tied economically through the token.

For miners, a rising token price increases the value of their reward pool. As profits grow, competitive miners reinvest in their operations, which improves the Subnet's collective output and forces other miners to keep pace.

For Subnet Owners, that improvement is critical. The products being built on top of a Subnet only get better if miners keep improving, so there's a direct incentive to drive demand into the token. Revenue-driven buybacks strengthen miner incentives, which leads to better outputs, which makes the product easier to monetize.

When buybacks exceed emissions, supply contracts and the token becomes deflationary. This deflation drives the flywheel that makes Subnet economics durable: higher prices strengthen miner incentives, attracting more talent and better performance.

This creates a structural mechanism where product success directly funds the miners those products depend on.





## Putting Numbers Behind Leading Subnets

A Subnet token reflects the economic output of the system it coordinates. Its value ultimately comes from the scale and durability of the revenue the subnet generates.

Some projects choose to route a portion of that value back to token holders through buybacks to reinforce demand for the token and signal that the token is a central part of the project's design. But they're not the only path. Strong, growing revenue is what anchors long-term value, and buybacks are simply one mechanism teams can use to express that strength.

**That means the two most important drivers for a Subnet are:**

- 1. Revenue potential** – The team's ability to monetize miner outputs, the pace of revenue growth, and the scale of the market they're targeting.
- 2. Capital efficiency** – How efficiently the Subnet generates revenue relative to miner emission costs.

This is why Subnets that look similar on paper can behave very differently. Two Subnets incentivizing miners to provide GPUs might each use that compute for entirely different workloads, target different users, and follow different go-to market strategies. Even with identical miner outputs, the economic dynamics can diverge sharply. Each Subnet becomes its own economic system with its own demand profile and risk-return characteristics.



## Why FDV Isn't The Best Valuation Framework For Subnets

With the mechanics in place, the next step is choosing a valuation framework that fits Subnets, and that's not a fully-diluted valuation (FDV). FDV has been the default metric across broader crypto, but Subnet tokens are structurally different and require a different lens.

Subnet tokens launch directly into public markets with no pre-sales, insider allocations, or preferential pricing. Risk, upside, and emissions are distributed from day one. Each follows a transparent, Bitcoin-like emissions curve that starts high and gradually tapers over time. This structure calls for its own valuation approach, drawing from work that Arca pioneered in this area.

While Subnet token emissions are roughly twice as fast as Bitcoin's and there are variables that make exact forecasts difficult, the circulating supply trajectory remains reasonably predictable. After the TAO halving in December 2025, we expect new Subnets to follow roughly this schedule:

- ~18% (3.7M) by month 12; and
- ~36% (7.5M) by month 24.

For Subnets that have been live since the start of Dynamic TAO, we expect the next one and two years to track a similar pattern:

- ~36% (7.5M) by Dec. 2026; and
- ~51% (10.6M) by Dec. 2027.

This level of predictability makes headline FDV unnecessary as a valuation anchor. Instead, valuations should be tied to an investor's time horizon and the projected circulating supply at that point. Using that supply estimate, the FDV can be discounted accordingly. For example, if 36% of the supply is expected to be circulating, the appropriate discount is 64%.

Looking at Subnets through that lens, Subnets end up looking far more attractively priced than their headline FDVs suggest.



## Let's Run Through Some Examples

### Chutes



Chutes has found clear product-market fit and is currently one of the top open source inference providers on OpenRouter. It provides over 110B tokens per day and is earning \$5.8M in annualized revenue.

With the AI inference market projected to reach \$255B by 2030, Chutes' growth potential appears strong. Fireworks AI, a comparable company also serving open-source models, 20x'd its revenue year-over-year and recently raised at a 31x multiple. If Chutes can 10x its current revenue by the end of 2026 and trades at a 30x multiple, it could clear a \$1.5B valuation.

On a one-year horizon, with ~36% of the supply circulating, a 64% discount puts its effective valuation at roughly \$174M—suggesting a 7x upside.

### Ridges



Ridges is in what is likely the hottest sector of the AI application space - AI coding agents. AI coding agent usage, revenue, and valuations have exploded in 2024-2025. Cursor recently announced a \$2.3B raise at a \$29.3B valuation, going from \$0 to \$100M revenue in its first year, and recently sitting at \$500M in annualized revenue. Claude Code was introduced in February 2025 by Anthropic and is pacing to hit \$9B in ARR by the end of the year.

The speed from product release to 9-figure annualized revenue figures is astounding and speaks to the market appetite for coding agents. Startups in the space earn high 20x-40x multiples based on the pace of revenue growth and developer demand: Cursor: 20x, Windsurf: 30x, Lovable: 40x.

Using a 30x multiple and 85,000 expected users at \$29 per month, Ridges could reach a ~\$720 million valuation. Ridges currently trades at a \$268M FDV. Adjusting by 64% for a one-year horizon from today yields an effective \$91M valuation—suggesting an 8x upside.



## Let's Run Through Some Examples

### Targon



The Targon team is ruthlessly focused on growing its top line by targeting enterprise customers.

The team is already earning \$10M in annualized revenue, and we expect the team to surpass \$40M by the end of next year. With a 25x multiple, Targon's valuation could reach \$1B.

Targon currently trades at a \$268M FDV. Adjusting its FDV by 64% for a one-year horizon from today implies a \$91M effective valuation—suggesting an 11x opportunity.

### Lium



Since July, Lium has scaled to a \$5.3M annualized revenue run rate while sustaining ~30% rental rates. Its single-command deployment and renter dashboard make it the easiest and lowest-cost compute marketplace in the space.

We expect Lium's revenue to grow from \$5.3M today to over \$30M by the end of 2026. Compute infrastructure startups often command 25–35x multiples; at 25x, Lium would reach a ~\$750M valuation.

Lium currently trades at a \$350M FDV. Adjusting by 64% for a one-year horizon from today gives an effective \$119M valuation—suggesting a 6x opportunity.



## A Valuation Model For TAO

TAO's value stems from the collective value of all the Subnets on Bittensor. It's similar to how Y Combinator would be valued as the sum of the stakes it holds in its portfolio companies, plus a premium for the future startups that will join its incubator.

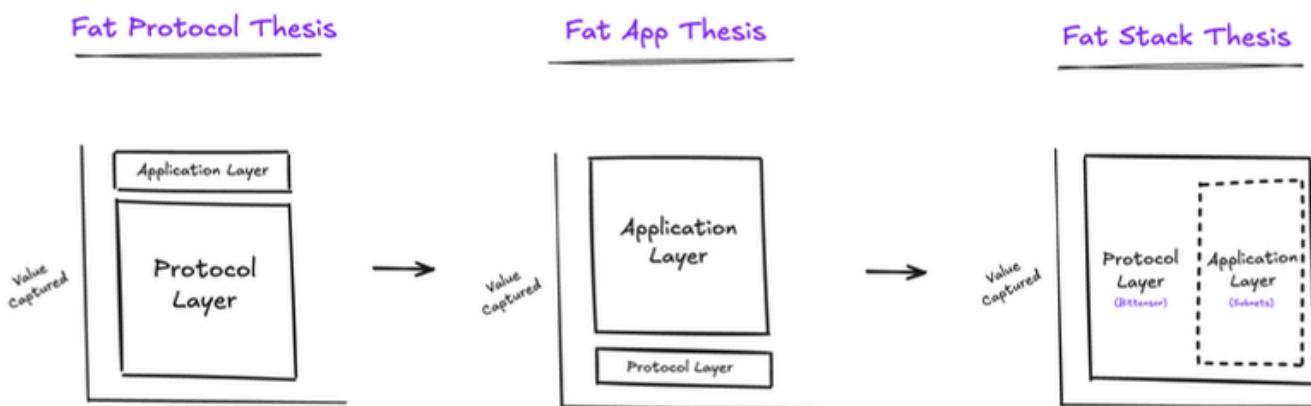
## The Fat Stack Thesis

The key to understanding our sum-of-the-parts methodology and why it's unique to Bittensor is the protocol's enshrined decentralized exchange (DEX). Each Subnet has its own DEX, which mandates that all Subnet tokens in the network are paired with TAO liquidity in a Uniswap-style AMM pool. TAO is the only medium of exchange to access Subnet tokens. Even as infrastructure matures and centralized exchanges start listing Subnet tokens, constant TAO emissions into the native DEX ensure that it will always be the deepest liquidity pool and remain the nexus of Subnet price discovery.

This design creates a direct and unavoidable link between Subnet growth and TAO demand.

The crypto ecosystem has long debated how value is captured between apps and the underlying protocol—in this case, Subnet tokens and TAO. We've been lectured at length about the relative value capture of tokens – from the [Fat Protocol Thesis](#) to the [Fat App Thesis](#)—characterizing L1s and the applications built on top of them as locked in a zero-sum battle for value accrual.

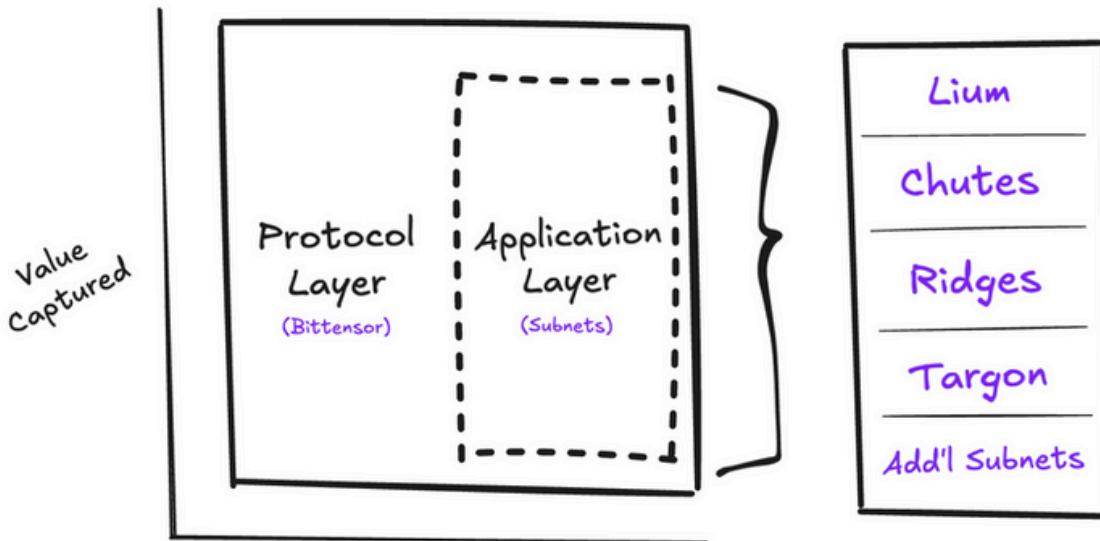
Bittensor introduces a third paradigm and deserves its own framework. We'll call this **The Fat Stack Thesis**. Rather than choosing between protocol or application, Bittensor's architecture captures value from both, simultaneously. When Subnets grow, they attribute higher value to TAO as the sole liquidity provider. When the Bittensor protocol grows, TAO gives Subnets and their miners even greater incentives. This symbiotic relationship, enforced by Bittensor's enshrined DEX that fixes Subnet token liquidity to TAO, defines the Fat Stack Thesis.





The applications (Subnets) on Bittensor grow the value of the network and that value is shared with the protocol (TAO). This isn't the case with L1s like Ethereum or Solana, where applications routinely extract value from the L1 rails and return less by circumventing base transaction fees and using alternative currencies (e.g. stablecoins) for liquidity. TAO functions as an index fund, explicitly tied to the collective value of its Subnets which can only be denominated in TAO.

## Sum-of-the-Parts Methodology



Our valuation framework builds on the enshrined DEX mechanism to establish an intrinsic, liquidity-driven price for TAO. We lean on a sum-of-the-parts (SOTP) methodology and take a bottoms-up approach, valuing the top Subnets based on private market comps.

UC Base Case			
	Dec-2026	Dec-2027	Dec-2028
<b>Total Subnets</b>	128	180	232
New Subnets		52	52
<b>Top Subnets</b>	6	9	12
Subnets representing 80% of network value	5%	5%	5%
<b>Top Subnets Marketcap (\$)</b>	\$5,760M	\$24,300M	\$93,960M
Avg. Valuation of Top Subnet	\$900M	\$2,700M	\$8,100M
Avg. Top Subnet ARR	\$30M	\$90M	\$270M
Avg. ARR Multiple	30x	30x	30x
<b>Total Subnets Marketcap (\$)</b>	\$7,200M	\$30,375M	\$117,450M
Add'l Value Outside of Top Subnets	20%	20%	20%
<b>Liquidity Pool Value (\$)</b>	\$3,240M	\$12,150M	\$49,329M
Liq : Mcap Ratio	0.45	0.40	0.42
<b>TAO in Liquidity Pool</b>	2,345,600	3,256,000	4,296,000
Circulating TAO	11,728,000	13,024,000	14,320,000
% Staked	20%	25%	30%
<b>DEX Liquidity Asset Value</b>	\$1,381	\$3,732	\$11,483

Note: You can find our model, valuation assumptions, and sensitivities [here](#)



## Our Model Works As Such:

**1. Subnet Valuation:** We make assumptions around the valuation trajectories and annualized revenue multiples of the top 5% of Subnets to project the aggregate market capitalization of the Subnet ecosystem. We expect this small percentage at the top to represent roughly 80% of the total network value. Today, we see roughly this level of value distribution based on annualized revenue.

**2. Liquidity Pool Modeling:** Using each Subnet's liquidity pool as the foundation, we model the total dollar value of TAO liquidity required to support the aggregate Subnet market capitalization.

**3. TAO Supply in Subnets:** We project the share of total TAO supply staked or locked in Subnet liquidity pools, informed by the emission schedule, historical TAO migration from the Root network, and behavioral patterns among miners, validators, and Subnet Owners.

**4. TAO Price Derivation:** Combining the total value of TAO locked in Subnet liquidity pools with the projected quantity of TAO staked across the ecosystem, we derive a price for TAO in USD terms.

DEX Liquidity Asset Value	\$1,381	\$3,732	\$11,483
L1 Premium	\$688	\$1,123	\$1,558
TAO Price	\$2,069	\$4,854	\$13,040

Finally, since TAO also derives value from its role as the store of value and medium of exchange within the network, it warrants an additional 'L1 premium.' L1 premiums are not derived from an onchain metric —for reference, Solana trades ~200x multiple to protocol fees. The premium is an underlying trust in the network and the asset as a non-sovereign store of value.

TAO is the trusted currency within Bittensor, similar to SOL and ETH within their respective networks. These networks are independent economies, bootstrapping a central store of value, defended by the network's activity. Native ecosystem tokens trade relative to alternative stores of value (e.g., SOL/ETH, ETH/BTC). Using this rationale and assuming a percent share of ETH and SOL market cap, we apply a conservative L1 premium to the economic value in the network.

## December 2027

<u>Scenario</u>	<u>Dec-2027</u>	<u>Dec-2028</u>
Bull	\$10,844	\$38,941
<b>Base</b>	<b>\$4,854</b>	<b>\$13,040</b>
Bear	\$663	\$1,361

As mentioned, several Subnets are already generating between \$5–10 million in annualized revenue today. By the end of 2027, we estimate that the leading Subnets could on average reach \$90M in annualized revenue—a 10-18x increase from current levels. Given the growth trajectories already observed across the top networks, we believe this range is achievable over a 24-month time horizon.

Applying a 30x revenue multiple (consistent with high growth SaaS and AI companies in the current market environment) yields an aggregate valuation of approximately \$24B for the top 5% of Subnets. Assuming the remaining 95% of Subnets collectively represent 20% of that value brings the total Subnet market capitalization to \$30B.

Using our SOTP model, we then apply the ratio of Subnet market capitalization to TAO liquidity pool depth and the projected share of TAO staked in the Subnet ecosystem. This produces an implied base-case price of \$3,732 per TAO.

Adding an L1 premium—to account for TAO’s role as the store of value within the network—suggests a base-case target of \$4,854 per TAO by December 2027.

Our bull case represents a scenario where the advantages of crypto-AI are widely recognized. As more Subnets prove they are a superior way to organize and operate the next generation of AI startups, the category sheds its “second-class” stigma. In this scenario, we model the top Subnets to reach approximately \$160M in annualized revenue on average by 2027. With crypto-AI no longer viewed as experimental, investor appetite expands and valuation multiples could reach higher. Our bull scenario suggests TAO could trade at an implied price of \$10,844.

In the bear case, we forecast a scenario where centralized AI gets so deeply embedded in the digital fabric of our lives that alternative products, even with higher performance, are slow to gain adoption. Passive users are satisfied by ‘good enough’ products, and inertia effectively outweighs the benefits of something technically superior. In this environment, traditional AI startups also face pressure. Our bear case models the shift to open-source competition taking much longer than anticipated, with TAO potentially trading at \$663 in 2027.

For now, that's not the world we live in. The AI landscape is in a state of hyper-competition, a relentless chase for performance where developers will swap out any product that falls behind.



## Caveats

### 1. Value Capture Assumption

We assume that Subnets will use their tokens as the primary capital asset whose value gets driven to vs others like equity. We've seen leading Subnets take a token-only approach to their business, though a minority of Subnets have opted to raise equity in addition to a Subnet token. In a token-only model, everything is designed to drive value back into the token. Once equity is introduced, that alignment becomes less direct. The team now has two assets to manage, which can create tension over where value should go and introduces the risk that Subnet Owners will choose to direct disproportionate value to the equityholders over the tokenholders.

Even with equity in place, Subnet Owners still have a structural reason to support the token since the product is deeply dependent on miners for high-quality outputs. At minimum, Subnet Owners need to recycle enough capital into buybacks to offset emissions.

A few teams are experimenting with legal frameworks to formally link the two, aligning token holders and shareholders under the same economic logic. But today, most Subnets continue to operate purely through tokens, which remains the simpler and cleaner model.

### 2. Subnets Staying on Bittensor

We assume successful Subnets will choose to stay on Bittensor. The main reason is that miners—the heart of every Subnet—may not follow them if they leave. Moving to another chain or to a standalone network creates uncertainty, introduces new trust assumptions, and could come with weaker incentives and shallower liquidity. TAO has the deepest liquidity pool in AI, and miners prefer ecosystems where they can convert earnings into stable value without friction.

Beyond the economics, Subnets benefit from the broader network. Bittensor is a positive-sum environment: the success of one Subnet strengthens TAO, and that rising TAO supports every other Subnet. Leaving means walking away from shared growth.

There are also emissions. TAO emissions function as recurring funding and a competitive moat. Giving that up would be irrational for any Subnet that is already succeeding inside the ecosystem. Taken together, these factors create a strong incentive for successful Subnets to stay on Bittensor rather than operate elsewhere.

# Outlook



## Looking into 2026

Bittensor will enter 2026 with two bullish network catalysts. First, the Emission Distribution Change, implemented in Nov. 2025, shifted the protocol's focus from price to net flows into AMM pools as the signal for distributing emissions. Subnets must now demonstrate positive net flows to receive emissions, effectively requiring real economic activity rather than speculative interest to create demand for their tokens. This forces capital efficiency across the ecosystem, ensuring TAO inflation isn't wasted and instead flows towards Subnets producing real value.

Second, the December 2025 Halving will reduce daily TAO emissions from 7,200 to 3,600, cutting new supply by 50% while Bittensor's economic activity continues expanding. Historical precedent from Bitcoin's halvings suggests this supply shock could positively and meaningfully impact price dynamics.

Together, these changes create a more disciplined capital allocation framework. Underperforming Subnets face economic pressure to improve or exit, while successful ones capture an increasing share of a shrinking emission pool.

## Hitting Escape Velocity

It's only been ten months since Dynamic TAO went live, yet the network already feels different. Development velocity has accelerated across every corner of the ecosystem. Subnets are pushing toward state-of-the-art results, telling clearer stories, generating revenue, building communities, and experimenting with token mechanics faster than at any point in Bittensor's history.

This level of progress within Bittensor stands out even more when you look at the rest of the crypto-AI category. We are researchers first, and we've looked at all the competitors. All remain theoretical or subscale by comparison. There are simply no serious contenders attempting what Bittensor has already proven at scale. And if a new one does emerge, it will face a structural disadvantage from day one, because it would need to replicate Bittensor's fair launch, one of its greatest moats. Pulling off a true fair launch today is nearly impossible. They only work when a project begins quietly and allows supply, community, and price to form organically over time.

Bittensor's fair launch created a foundation strong enough to withstand any market cycle. Even if the network were left untouched for years, it would keep running, and someone, somewhere, would eventually return to build on it.

We often think about Bittensor's success through a venture lens. Not every Subnet will work, nor does it need to. Like any incubator or venture portfolio, a handful of breakout successes can carry the entire portfolio. By traditional venture math, one in thirty makes the fund. Applied here, four or five successful Subnets could make Bittensor a generational success. Today, we already see more than ten approaching product-market fit.

If we're right, and early Subnet performance suggests we are, then TAO at current levels represents the kind of structural mispricing that may define generational investment opportunities.

# Disclaimers

This presentation is for informational purposes only.

It does not constitute investment advice or a recommendation to buy or sell any security.

Investments involve risk and there is no guarantee of return.

Past performance is not indicative of future results.

Consult your financial advisor before making any investment decisions.

The information contained herein is subject to change without notice.

Neither the author nor the publisher accept responsibility for any inaccuracies.

Any links to external websites are provided for convenience only and are not affiliated with the author or publisher.

The author and publisher disclaim any liability for damages resulting from the use of this information.

© 2023 [Author/Publisher]. All rights reserved.



# Disclaimers

## Private Fund Disclosure

Unsupervised Capital manages a private investment fund (the “Fund”) that qualifies for exclusion from the definition of investment company under Section 3(c)(1) of the Investment Company Act of 1940. The Fund is not registered under the Investment Company Act of 1940 or the Securities Act of 1933.

## This Report Is Not a Fund Offering

This report does not constitute an offer to sell or a solicitation of an offer to buy interests in the Fund or any other investment vehicle managed by Unsupervised Capital. The Fund is not available to the general public and is offered only to a limited number of qualified investors through separate, private channels in accordance with applicable securities laws. Investment in the Fund is restricted to accredited investors (as defined in Rule 501(a) of Regulation D) who meet additional suitability requirements and is subject to significant investment minimums, lock-up periods, and restrictions on redemption.

## No General Solicitation

Unsupervised Capital is not soliciting investors through this report. This report is provided solely for informational and educational purposes regarding our market views on Bittensor, digital assets, and related technologies. It is not intended to market, promote, or solicit investment in the Fund or any other vehicle we manage.

## Fund Performance Not Disclosed

This report does not contain any information regarding the Fund's performance, investment positions, or portfolio holdings. Any investment theses, positions, or strategies discussed reflect our general market views and do not necessarily represent the Fund's actual portfolio or investment decisions. The Fund's actual holdings, performance, and strategies may differ materially from any views expressed in this report.

## Investment Adviser Registration

Unsupervised Capital is registered with the U.S. Securities and Exchange Commission as an investment adviser under the Investment Advisers Act of 1940 (SEC File Number: 805-5052713762). Our Form ADV Part 2A brochure contains important information about our business practices, services, fees, and conflicts of interest.

Our Form ADV Part 2A brochure is available:

- At <https://reports.adviserinfo.sec.gov/reports/ADV/335890/PDF/335890.pdf>;
- Or upon request by submission [here](#).

Registration with the SEC does not imply any level of skill, training, or endorsement by the SEC.

## Existing Fund Investors

Current investors in the Fund should refer to the Fund's offering documents, limited partnership agreement, and quarterly reports for information regarding the Fund's operations, performance, and holdings. This report is not a substitute for those documents and does not constitute investor communications required under the Fund's governing documents.

## Not Investment Advice

This report constitutes general information and is not investment advice, financial advice, legal advice, or a recommendation to buy, hold, or sell any security or digital asset, including TAO or Subnet tokens. Receipt or review of this report does not create an investment advisory relationship, client relationship, or any other professional relationship between you and Unsupervised Capital. We are not acting as your adviser, fiduciary, or agent. This report does not constitute personalized investment advice tailored to your individual circumstances. Do not use this report as the sole basis for investment decisions. Consider all relevant information, including your existing portfolio, investment objectives, risk tolerance, liquidity needs, and investment time horizon. Consult qualified financial, legal, and tax advisors before making investment decisions.



## Disclaimers Cont.

### **Material Risks**

Digital assets, including TAO and Subnet tokens, are highly speculative and volatile, involving substantial risk of loss including potential total loss of capital. Specific risks include extreme price volatility, liquidity constraints, regulatory uncertainty, technology failures, security breaches, competitive obsolescence, custody risks, and lack of investor protections. The regulatory status of digital assets is uncertain and evolving; regulatory actions could materially and adversely affect their value, legality, or utility. Digital assets may be deemed securities under applicable laws, subjecting them to registration requirements and trading restrictions.

### **Forward-Looking Statements and Assumptions**

This report contains forward-looking statements, projections, and estimates regarding future events, valuations, revenue, and token prices. These statements are based on assumptions that may prove incorrect. Actual results may differ materially and adversely from projections. Factors causing differences include regulatory changes, technological failures, competitive developments, execution risk, and market volatility. Information is current as of December 2025 and may quickly become outdated. We undertake no obligation to update this report.

### **Valuation Model Limitations**

Price targets and valuation analyses in this Report are based on a public model that employs subjective assumptions and methodologies. Different assumptions or methodologies would produce different results. The valuation model has inherent limitations and does not purport to represent the only possible approach or outcome. Specific assumptions regarding revenue growth, market multiples, liquidity ratios, staking behavior, and other inputs are speculative and may prove materially incorrect. Sensitivity to assumption changes may be significant. For additional discussion of model limitations, see the separate Valuation Model Disclaimer.

### **Conflicts of Interest**

Unsupervised Capital and its affiliates, employees, and associated persons hold positions in TAO and various Subnet tokens. We have material financial interests in the appreciation of TAO and Bittensor network success. This creates inherent conflicts of interest and potential bias in our analysis. We may buy, sell, trade, stake, or participate in network activities involving TAO or Subnet tokens at any time without notice. We may provide advisory services or investment to Subnet teams and may have relationships with entities affiliated with Bittensor or specific Subnets. Unsupervised Capital does not owe any fiduciary duty to readers of this report. We have no obligation to act in your best interest or to disclose conflicts beyond those stated herein.

In addition to holding positions in TAO and Subnet tokens, the following conflicts may exist:

- We may receive indirect economic benefits from TAO price appreciation beyond our direct holdings through fund management fees calculated on asset value
- We may have relationships with Bittensor co-founders, Subnet teams, developers, or other Bittensor participants that create additional conflicts
- Our analysis may be influenced by our desire to support the market value of our existing positions

These conflicts may create incentive to present information in a manner favorable to our financial interests.

### **Information Sources and Verification**

Information is derived from sources believed reliable, including publicly available data, Subnet team disclosures, and our own analysis. This report does not contain material non-public information. We have not independently verified all information and make no representation regarding its accuracy, completeness, or reliability. Revenue figures, user metrics, and performance data may be self-reported, unaudited, based on limited time periods, or subject to different methodologies. You are solely responsible for conducting independent due diligence and verification.

### **No Warranty; Limitation of Liability**

THIS REPORT IS PROVIDED "AS IS" WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED. TO THE MAXIMUM EXTENT PERMITTED BY LAW, UNSUPERVISED CAPITAL SHALL NOT BE LIABLE FOR ANY DAMAGES ARISING OUT OF OR RELATED TO YOUR USE OF OR RELIANCE ON THIS REPORT, INCLUDING DIRECT, INDIRECT, INCIDENTAL, CONSEQUENTIAL, OR PUNITIVE DAMAGES, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.



## Disclaimers Cont.

### Intellectual Property

This report is protected by copyright and intellectual property laws. You may not reproduce, distribute, modify, create derivative works, or publicly display this report without prior written permission, except for personal, non-commercial use as permitted by law. Unauthorized use may violate copyright, trademark, and other laws.

### Third-Party Content

This report may reference or link to third-party websites, services, or content. We do not endorse, control, or assume responsibility for third-party content. Links are provided for convenience only. You access third-party content at your own risk.

### Jurisdictional Restrictions

This report is not intended for distribution in any jurisdiction where such distribution would be contrary to law or require registration. Persons accessing this report are responsible for compliance with applicable laws in their jurisdiction.

#### Hypothetical Performance

Price targets, valuations, and projections in this report constitute hypothetical performance. Hypothetical performance has inherent limitations and does not represent actual trading or investment results.

#### Material Assumptions:

- Subnet revenue growth of 10-18x over 24 months
- Valuation multiples of 20x-40x revenue consistent with high-growth SaaS companies
- TAO liquidity ratios and staking percentages based on early network behavior
- Continuation of current regulatory environment and market conditions
- Subnet token buyback mechanisms functioning as designed

#### Limitations:

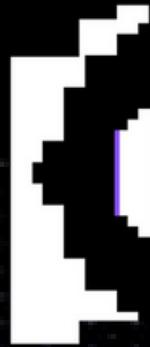
- Assumptions may not materialize; actual results may differ materially
- Does not reflect impact of fees, expenses, or taxes
- Benefits from hindsight and ability to select favorable assumptions
- Does not account for all risks or market conditions
- Model sensitivity: small changes in assumptions produce significantly different outcomes
- No representation that any investment will achieve similar results

The projections contained in this report do not reflect actual outcomes and should not be relied upon as indicative of future results.

### Selective Presentation of Investments

This report discusses certain Subnets that represent favorable examples of the Bittensor ecosystem. This presentation is selective and does not reflect all investments, positions, or market views held by Unsupervised Capital. We may not have discussed underperforming or failed Subnets or Subnets that have been deregistered from the network.

The Subnets discussed were selected to illustrate specific points regarding the Bittensor model and do not represent all available investment opportunities within the network. The success of highlighted Subnets does not predict success of other Subnets.



UNSUPERVISED  
CAPITAL

UNSUPERVISED.CAPITAL

Contact us at [contact@unsupervised.capital](mailto:contact@unsupervised.capital)